

LAB5 Logistic Regression Part3&4

March 4, 2021

1 LAB5 KNN Classification on Diabetes Part3

In this part I will use the diabetes dataset and create a KNN classifier. Tune the parameter K to find the optimal K for the classification task. Then perform k-fold cross validation, find the best k that minimizes the misclassification rate.

```
[1]: import pandas as pd
import numpy as np
import matplotlib as plt
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
from sklearn.datasets import load_breast_cancer
from sklearn import linear_model
from mpl_toolkits.mplot3d import Axes3D
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc
import seaborn as sns
import matplotlib.patches as mpatches
import operator
from matplotlib.lines import Line2D
from sklearn.model_selection import cross_val_score
```

```
[2]: df = pd.read_csv('diabetes.csv')
```

1.0.1 Data Manuplation

```
[3]: #diabetes.columns = ['x1', 'x2', 'x3', 'x4', 'x5', 'x6', 'x7', 'x8', 'k']
df.columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
```

```
[4]: df['c'].replace({0: 'A', 1: 'B'}, inplace=True)
df
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	1	85	66	29	0	26.6	
1	8	183	64	0	0	23.3	
2	1	89	66	23	94	28.1	

3	0	137	40	35	168	43.1
4	5	116	74	0	0	25.6
..
762	10	101	76	48	180	32.9
763	2	122	70	27	0	36.8
764	5	121	72	23	112	26.2
765	1	126	60	0	0	30.1
766	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	c
0	0.351	31	A
1	0.672	32	B
2	0.167	21	A
3	2.288	33	B
4	0.201	30	A
..
762	0.171	63	A
763	0.340	27	A
764	0.245	30	A
765	0.349	47	B
766	0.315	23	A

[767 rows x 9 columns]

Pregnancies: x1
 Glucose: x2
 BloodPressure: x3
 SkinThickness: x4
 Insulin: x5
 BMI: x6
 DiabetesPedigreeFunction: x7
 Age: x8
 Outcome: k (class)

[]:

[5]: df.c

[5]: 0 A
 1 B
 2 A
 3 B
 4 A
 ..
 762 A
 763 A
 764 A

```
765    B
766    A
Name: c, Length: 767, dtype: object
```

```
[6]: df_array = df.to_numpy()
df_array
```

```
[7]: array([[1, 85, 66, ..., 0.35100000000000003, 31, 'A'],
           [8, 183, 64, ..., 0.672, 32, 'B'],
           [1, 89, 66, ..., 0.16699999999999998, 21, 'A'],
           ...,
           [5, 121, 72, ..., 0.245, 30, 'A'],
           [1, 126, 60, ..., 0.349, 47, 'B'],
           [1, 93, 70, ..., 0.315, 23, 'A']], dtype=object)
```


```
[7]: df_array[:, :8] #all features except c, select all columns
```

```
[7]: array([[1, 85, 66, ..., 26.6, 0.35100000000000003, 31],
           [8, 183, 64, ..., 23.3, 0.672, 32],
           [1, 89, 66, ..., 28.1, 0.16699999999999998, 21],
           ...,
           [5, 121, 72, ..., 26.2, 0.245, 30],
           [1, 126, 60, ..., 30.1, 0.349, 47],
           [1, 93, 70, ..., 30.4, 0.315, 23]], dtype=object)
```

```
[8]: df_array[:, [0, 3]] #selecting specific features, ex: first feature and forth
      ↪ feature
```

```
[8]: array([[1, 29],
           [8, 0],
           [1, 23],
           ...,
           [5, 23],
           [1, 0],
           [1, 31]], dtype=object)
```

```
[9]: sns.
      ↪ lmpplot('Age', 'BMI', data=df, hue='c', palette='Set1', fit_reg=False, scatter_kws={"s":
      ↪ 30})
plt.show()
```



```
[10]: #--- To calculate the distance between two points---#
```

```
def euclidean_distance(pt1,pt2,dimension):
    distance=0
    for x in range(dimension):
        distance += np.square(pt1[x]-pt2[x])
    return np.sqrt(distance)
```

```
[11]: #---our own KNN model---#
```

```
def knn(training_points,test_point,k):
    distances={}

    #---the number of axes we are dealing with---#
    dimension = test_point.shape[1]

    #---calculating euclidean distance between each point
    # in the training data and test data
    for x in range(len(training_points)):
        dist= euclidean_distance(test_point,training_points.iloc[x],dimension)
```

```

    #--- record the distance for each training points---#
    distances[x]=dist[0]

    #---sort the distances---#
    sorted_d=sorted(distances.items(),key=operator.itemgetter(1))

    #---to store the neighbors---#
    neighbors=[]

    #---extract the top k neighbors---#
    for x in range(k):
        neighbors.append(sorted_d[x][0])

    #---for each neighbor found, find its class---#
    class_counter={}
    for x in range(len(neighbors)):
        #---find out the class for that particular point---#
        cls=training_points.iloc[neighbors[x]][-1]

        if cls in class_counter:
            class_counter[cls] += 1
        else:
            class_counter[cls]=1

        #---sort the class_counter in descending order---#
        sorted_counter=sorted(class_counter.items(), key=operator.
→itemgetter(1),reverse=True)

    #---return the class with the most count, as well as the neighbors found---#
    return(sorted_counter[0][0],neighbors)

```

```

[12]: test_set=[[30,40]]                                #test_set[Age, BMI]
test=pd.DataFrame(test_set)

#Choose features specifically in dataframe for passing to knn func.
cls,neighbors=knn(df[["Age", "BMI", "c"]],test,3)
→#knn(training_points, test_points, k)
print("Predicted Class:" + cls)
#print(len(neighbors))

```

Predicted Class:A

```

[13]: #---Test Point---#
#test_set=[[32,23]]                                #test_set[Age, BMI]
#test=pd.DataFrame(test_set)
#cls,neighbors=knn(df,test,3)                        # This is an issue

```

```
#print("Predicted Class:" + cls)
```

Issue was my passing points in function was bit off.

1.0.2 Visualizing Different Values of K

It is useful to be able to visualize the effect of applying various values of k. The following code snippet draws a series of concentric circles around the test point based on the values of k, which range from 7 to 1, with intervals of -2:

```
[ ]:
```

```
[14]: df = pd.read_csv('diabetes.csv')
df.columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',
               ↪ 'Diabetes']
df_array = df.to_numpy()
df_array[:, :8]
```

```
[14]: array([[1.00e+00, 8.50e+01, 6.60e+01, ..., 2.66e+01, 3.51e-01, 3.10e+01],
             [8.00e+00, 1.83e+02, 6.40e+01, ..., 2.33e+01, 6.72e-01, 3.20e+01],
             [1.00e+00, 8.90e+01, 6.60e+01, ..., 2.81e+01, 1.67e-01, 2.10e+01],
             ...,
             [5.00e+00, 1.21e+02, 7.20e+01, ..., 2.62e+01, 2.45e-01, 3.00e+01],
             [1.00e+00, 1.26e+02, 6.00e+01, ..., 3.01e+01, 3.49e-01, 4.70e+01],
             [1.00e+00, 9.30e+01, 7.00e+01, ..., 3.04e+01, 3.15e-01, 2.30e+01]])
```

```
[15]: #display(X)
```

```
[16]: #display(y)
```

```
[17]: %matplotlib inline

datasets = df_array
X=df_array[:, [5,7]]          ##### take the "Age" and "BMI" features
y=df.c # k = target (classes)

colors=['red', 'green',]

for color, i, k in zip(colors, [0,1], df.c):
    plt.scatter(X[y==i,0], X[y==i,1], color=color, label=i+1)

plt.xlabel('Age')
plt.ylabel('BMI')


#custom = [Line2D([], [], marker='.', color='r', linestyle='None'), Line2D([], [],
    ↪ [], marker='.', color='g', linestyle='None')]
```

```

plt.legend(custom, ['Yes', 'No'], loc='lower right')
plt.legend(["1:Yes Diabetes","2:No Diabetes"], loc='best', shadow=False,
↳scatterpoints=1)

plt.title("Diabetes Classifier")
plt.show()

```



```

[18]: from sklearn.neighbors import KNeighborsClassifier

k=5
#---instantiate learning model---#
knn=KNeighborsClassifier(n_neighbors=k)
#---fitting the model---#
knn.fit(X,y)

#---min and max for the first feature---#
x_min, x_max=X[:,0].min()-1,X[:,0].max() + 1

#---min and max for the second feature---#
y_min, y_max=X[:,1].min()-1, X[:,1].max() + 1

#---step size in mesh---#
h=(x_max/x_min)/100

```

```

#---make predictions for each of the points in xx,yy---#
xx,yy=np.meshgrid(np.arange(x_min,x_max,abs(h)),
                  np.arange(y_min,y_max,abs(h)))

Z=knn.predict(np.c_[xx.ravel(),yy.ravel()])

#---draw the result using a color plot---#
Z=Z.reshape(xx.shape)
plt.contourf(xx,yy,Z,cmap=plt.cm.Accent,alpha=0.5)

#---plot the training points---#
colors=['red','blue']
for color, i, k in zip(colors,[0,1],df.c):
    plt.scatter(X[y==i,0],X[y==i,1],color=color,label=i+1)

plt.xlabel('Age')
plt.ylabel('BMI')

plt.title(f'KNN (k=5)')

plt.legend(["Red: Benign Diabetes","Blue: Malignant Diabetes"], loc='best',
           shadow=False, scatterpoints=1)

predictions=knn.predict(X)


#---classifications based on predictions---#
print(np.unique(predictions,return_counts=True))

```

```

(array([0, 1]), array([515, 252]))

```

```
[19]: from sklearn.neighbors import KNeighborsClassifier

k=17
#---instantiate learning model---#
knn=KNeighborsClassifier(n_neighbors=k)
#---fitting the model---#
knn.fit(X,y)

#---min and max for the first feature---#
x_min, x_max=X[:,0].min()-1,X[:,0].max() + 1

#---min and max for the second feature---#
y_min, y_max=X[:,1].min()-1, X[:,1].max() + 1

#---step size in mesh---#
h=(x_max/x_min)/100

#---make predictions for each of the points in xx,yy---#
xx,yy=np.meshgrid(np.arange(x_min,x_max,abs(h)),
                  np.arange(y_min,y_max,abs(h)))

Z=knn.predict(np.c_[xx.ravel(),yy.ravel()])
```

```

#---draw the result using a color plot---#
Z=Z.reshape(xx.shape)
plt.contourf(xx,yy,Z,cmap=plt.cm.Accent,alpha=0.5)

#---plot the training points---#
colors=['red','blue']
for color, i, k in zip(colors,[0,1],df.c):
    plt.scatter(X[y==i,0],X[y==i,1],color=color,label=i+1)

plt.xlabel('Age')
plt.ylabel('BMI')

plt.title(f'KNN (k=17)')

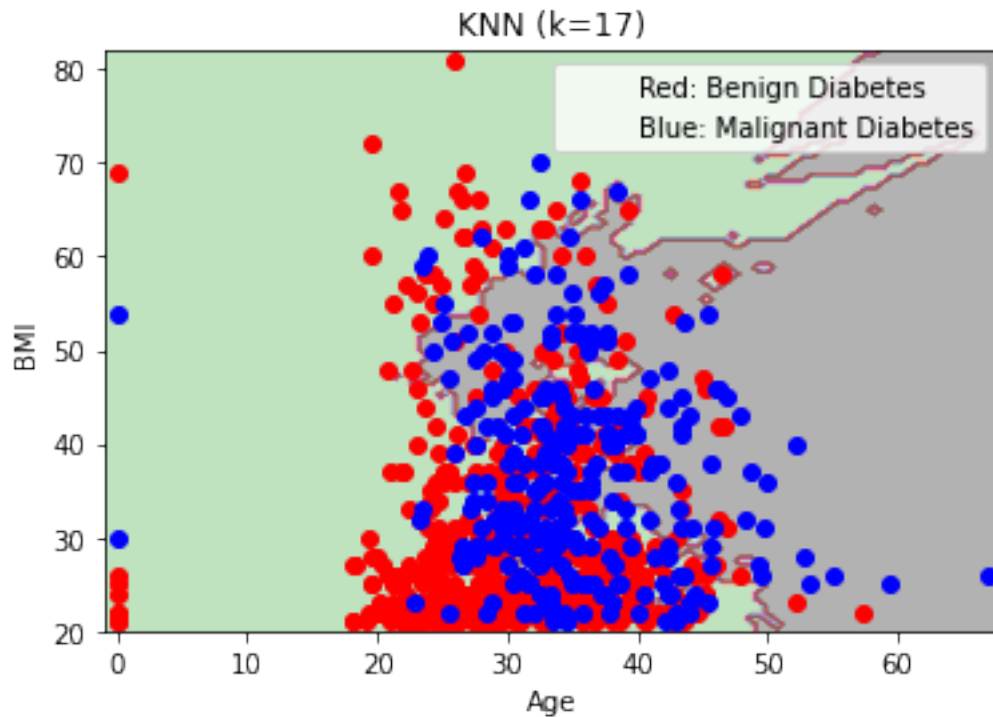
plt.legend(["Red: Benign Diabetes","Blue: Malignant Diabetes"], loc='best',
           shadow=False, scatterpoints=1)

predictions=knn.predict(X)

#---classifications based on predictions---#
print(np.unique(predictions,return_counts=True))

```

(array([0, 1]), array([525, 242]))



note that as k increases, the boundary becomes smoother. But it also means that the more points will be classified incorrectly. When k increases to a large value, underfitting occurs.

1.0.3 Cross-Validation

We have witnessed from previous chapters, that we split our dataset into two individual sets—one for training and one for testing. However, the data in your dataset may not be distributed evenly, and as a result your test set may be too simple or too hard to predict, thereby making it very difficult to know if your model works well.

Instead of using part of the data for training and part for testing, you can split the data into k -folds and train the model k times, rotating the training and testing sets. By doing so, each data point is now being used for training and testing.

Parameter-Tuning K Now that you understand cross-validation, let's use it on our Iris dataset. We will train the model using all of the three features, and at the same time we shall use cross-validation on the dataset using 10 folds. We will do this for each value of k :

```
[59]: #---holds the cv (cross-validates) scores---#
cv_scores=[]

#---use all features---#
X=df_array[:, :8]          ##### select all features except c, select all
    ↪ columns
#X=df_array[:, [5,7]]      ##### select Age and BMI
y=df.c # k = target (classes)

#---number of folds---#
folds=10

#---creating odd list of K for KNN---#
ks=list(range(1,int(len(X)*((folds-1)/folds))))

#---remove all multiples of 3---#
ks=[ck for ck in ks if ck % 3 !=0]

#---perform k-fold cross-validation---#
for k in ks:
    knn=KNeighborsClassifier(n_neighbors=k)

    #---performs cross-validation and returns the average accuracy---#
    scores= cross_val_score(knn,X,y,cv=folds,scoring='accuracy')
    mean=scores.mean()
    cv_scores.append(mean)
    print('K value: ', k, 'Average Accuracy: ', mean)
```

```
K value:  1 Average Accuracy:  0.6806390977443609
```

```
K value:  2 Average Accuracy:  0.7105775803144223
```

K value: 4 Average Accuracy: 0.7197368421052632
 K value: 5 Average Accuracy: 0.7249999999999999
 K value: 7 Average Accuracy: 0.7432843472317157
 K value: 8 Average Accuracy: 0.7380553656869446
 K value: 10 Average Accuracy: 0.7484278879015721
 K value: 11 Average Accuracy: 0.7380211893369788
 K value: 13 Average Accuracy: 0.7471291866028709
 K value: 14 Average Accuracy: 0.7549555707450445
 K value: 16 Average Accuracy: 0.7510765550239235
 K value: 17 Average Accuracy: 0.7550922761449078
 K value: 19 Average Accuracy: 0.7485645933014354
 K value: 20 Average Accuracy: 0.7524606971975393
 K value: 22 Average Accuracy: 0.7525119617224881
 K value: 23 Average Accuracy: 0.7499145591250855
 K value: 25 Average Accuracy: 0.7472488038277513
 K value: 26 Average Accuracy: 0.7342105263157894
 K value: 28 Average Accuracy: 0.7289644565960355
 K value: 29 Average Accuracy: 0.7328947368421053
 K value: 31 Average Accuracy: 0.7342105263157894
 K value: 32 Average Accuracy: 0.7342105263157894
 K value: 34 Average Accuracy: 0.7394053315105947
 K value: 35 Average Accuracy: 0.7420027341079973
 K value: 37 Average Accuracy: 0.7367908407382092
 K value: 38 Average Accuracy: 0.7433185235816815
 K value: 40 Average Accuracy: 0.7394395078605606
 K value: 41 Average Accuracy: 0.736825017088175
 K value: 43 Average Accuracy: 0.7316643882433356
 K value: 44 Average Accuracy: 0.7303827751196172
 K value: 46 Average Accuracy: 0.7251537935748462
 K value: 47 Average Accuracy: 0.7290840738209159
 K value: 49 Average Accuracy: 0.7316985645933014
 K value: 50 Average Accuracy: 0.723838004101162
 K value: 52 Average Accuracy: 0.7316473000683528
 K value: 53 Average Accuracy: 0.7329630895420369
 K value: 55 Average Accuracy: 0.7303485987696514
 K value: 56 Average Accuracy: 0.7303315105946686
 K value: 58 Average Accuracy: 0.7290328092959673
 K value: 59 Average Accuracy: 0.7316643882433357
 K value: 61 Average Accuracy: 0.7368079289131921
 K value: 62 Average Accuracy: 0.7355263157894737
 K value: 64 Average Accuracy: 0.7328947368421053
 K value: 65 Average Accuracy: 0.736825017088175
 K value: 67 Average Accuracy: 0.7394395078605606
 K value: 68 Average Accuracy: 0.7394395078605605
 K value: 70 Average Accuracy: 0.7367908407382091
 K value: 71 Average Accuracy: 0.7394224196855776
 K value: 73 Average Accuracy: 0.7368079289131921
 K value: 74 Average Accuracy: 0.7393711551606289

K value: 76 Average Accuracy: 0.7380895420369106
 K value: 77 Average Accuracy: 0.7368421052631579
 K value: 79 Average Accuracy: 0.7355092276144909
 K value: 80 Average Accuracy: 0.7289986329460014
 K value: 82 Average Accuracy: 0.7198564593301435
 K value: 83 Average Accuracy: 0.7289473684210526
 K value: 85 Average Accuracy: 0.7328605604921394
 K value: 86 Average Accuracy: 0.7315618591934382
 K value: 88 Average Accuracy: 0.730228981544771
 K value: 89 Average Accuracy: 0.7328434723171566
 K value: 91 Average Accuracy: 0.7237183868762816
 K value: 92 Average Accuracy: 0.7224196855775803
 K value: 94 Average Accuracy: 0.726332877648667
 K value: 95 Average Accuracy: 0.7224196855775803
 K value: 97 Average Accuracy: 0.7197881066302119
 K value: 98 Average Accuracy: 0.721103896103896
 K value: 100 Average Accuracy: 0.7198051948051948
 K value: 101 Average Accuracy: 0.7249999999999999
 K value: 103 Average Accuracy: 0.7249999999999999
 K value: 104 Average Accuracy: 0.7223855092276145
 K value: 106 Average Accuracy: 0.7132775119617224
 K value: 107 Average Accuracy: 0.7197881066302119
 K value: 109 Average Accuracy: 0.7119617224880382
 K value: 110 Average Accuracy: 0.7106630211893369
 K value: 112 Average Accuracy: 0.7145933014354068
 K value: 113 Average Accuracy: 0.7210868079289132
 K value: 115 Average Accuracy: 0.715892002734108
 K value: 116 Average Accuracy: 0.7158578263841421
 K value: 118 Average Accuracy: 0.7132775119617224
 K value: 119 Average Accuracy: 0.7184894053315106
 K value: 121 Average Accuracy: 0.7223684210526315
 K value: 122 Average Accuracy: 0.7223513328776486
 K value: 124 Average Accuracy: 0.7249658236500343
 K value: 125 Average Accuracy: 0.7249487354750512
 K value: 127 Average Accuracy: 0.7223684210526315
 K value: 128 Average Accuracy: 0.719771018455229
 K value: 130 Average Accuracy: 0.7158578263841421
 K value: 131 Average Accuracy: 0.7184723171565276
 K value: 133 Average Accuracy: 0.715892002734108
 K value: 134 Average Accuracy: 0.7119788106630212
 K value: 136 Average Accuracy: 0.7132775119617226
 K value: 137 Average Accuracy: 0.7159090909090909
 K value: 139 Average Accuracy: 0.7184894053315105
 K value: 140 Average Accuracy: 0.7145762132604239
 K value: 142 Average Accuracy: 0.7119788106630212
 K value: 143 Average Accuracy: 0.7184552289815447
 K value: 145 Average Accuracy: 0.7145420369104578
 K value: 146 Average Accuracy: 0.7066985645933015

K value: 148 Average Accuracy: 0.7041011619958988
 K value: 149 Average Accuracy: 0.7053998632946001
 K value: 151 Average Accuracy: 0.7066814764183186
 K value: 152 Average Accuracy: 0.7053827751196172
 K value: 154 Average Accuracy: 0.694958988380041
 K value: 155 Average Accuracy: 0.7028024606971975
 K value: 157 Average Accuracy: 0.7053998632946001
 K value: 158 Average Accuracy: 0.7027511961722488
 K value: 160 Average Accuracy: 0.7001708817498291
 K value: 161 Average Accuracy: 0.7066814764183185
 K value: 163 Average Accuracy: 0.7001367053998633
 K value: 164 Average Accuracy: 0.6923103212576897
 K value: 166 Average Accuracy: 0.6910116199589884
 K value: 167 Average Accuracy: 0.6975222146274779
 K value: 169 Average Accuracy: 0.6896958304853043
 K value: 170 Average Accuracy: 0.6831852358168147
 K value: 172 Average Accuracy: 0.6779733424470267
 K value: 173 Average Accuracy: 0.6805707450444294
 K value: 175 Average Accuracy: 0.680604921394395
 K value: 176 Average Accuracy: 0.6688824333561175
 K value: 178 Average Accuracy: 0.6740943267259057
 K value: 179 Average Accuracy: 0.6818865345181135
 K value: 181 Average Accuracy: 0.675393028024607
 K value: 182 Average Accuracy: 0.6688824333561175
 K value: 184 Average Accuracy: 0.6623718386876283
 K value: 185 Average Accuracy: 0.6649692412850308
 K value: 187 Average Accuracy: 0.6623376623376622
 K value: 188 Average Accuracy: 0.6584244702665755
 K value: 190 Average Accuracy: 0.6558270676691731
 K value: 191 Average Accuracy: 0.6544941900205059
 K value: 193 Average Accuracy: 0.6571257689678743
 K value: 194 Average Accuracy: 0.6519309637730691
 K value: 196 Average Accuracy: 0.6532296650717704
 K value: 197 Average Accuracy: 0.6545283663704717
 K value: 199 Average Accuracy: 0.6519309637730691
 K value: 200 Average Accuracy: 0.6532296650717704
 K value: 202 Average Accuracy: 0.6519309637730691
 K value: 203 Average Accuracy: 0.6519309637730691
 K value: 205 Average Accuracy: 0.6506151742993849
 K value: 206 Average Accuracy: 0.6519138755980862
 K value: 208 Average Accuracy: 0.6532125768967875
 K value: 209 Average Accuracy: 0.6532125768967875
 K value: 211 Average Accuracy: 0.6532125768967875
 K value: 212 Average Accuracy: 0.6532125768967875
 K value: 214 Average Accuracy: 0.6532125768967875
 K value: 215 Average Accuracy: 0.6532125768967875
 K value: 217 Average Accuracy: 0.6532125768967875
 K value: 218 Average Accuracy: 0.6519138755980862

[illegible]

[illegible]

[illegible]

[illegible]


```

K value: 652 Average Accuracy: 0.6519138755980862
K value: 653 Average Accuracy: 0.6519138755980862
K value: 655 Average Accuracy: 0.6519138755980862
K value: 656 Average Accuracy: 0.6519138755980862
K value: 658 Average Accuracy: 0.6519138755980862
K value: 659 Average Accuracy: 0.6519138755980862
K value: 661 Average Accuracy: 0.6519138755980862
K value: 662 Average Accuracy: 0.6519138755980862
K value: 664 Average Accuracy: 0.6519138755980862
K value: 665 Average Accuracy: 0.6519138755980862
K value: 667 Average Accuracy: 0.6519138755980862
K value: 668 Average Accuracy: 0.6519138755980862
K value: 670 Average Accuracy: 0.6519138755980862
K value: 671 Average Accuracy: 0.6519138755980862
K value: 673 Average Accuracy: 0.6519138755980862
K value: 674 Average Accuracy: 0.6519138755980862
K value: 676 Average Accuracy: 0.6519138755980862
K value: 677 Average Accuracy: 0.6519138755980862
K value: 679 Average Accuracy: 0.6519138755980862
K value: 680 Average Accuracy: 0.6519138755980862
K value: 682 Average Accuracy: 0.6519138755980862
K value: 683 Average Accuracy: 0.6519138755980862
K value: 685 Average Accuracy: 0.6519138755980862
K value: 686 Average Accuracy: 0.6519138755980862
K value: 688 Average Accuracy: 0.6519138755980862
K value: 689 Average Accuracy: 0.6519138755980862

```

```
[60]: print(ks)
```

```

[1, 2, 4, 5, 7, 8, 10, 11, 13, 14, 16, 17, 19, 20, 22, 23, 25, 26, 28, 29, 31,
32, 34, 35, 37, 38, 40, 41, 43, 44, 46, 47, 49, 50, 52, 53, 55, 56, 58, 59, 61,
62, 64, 65, 67, 68, 70, 71, 73, 74, 76, 77, 79, 80, 82, 83, 85, 86, 88, 89, 91,
92, 94, 95, 97, 98, 100, 101, 103, 104, 106, 107, 109, 110, 112, 113, 115, 116,
118, 119, 121, 122, 124, 125, 127, 128, 130, 131, 133, 134, 136, 137, 139, 140,
142, 143, 145, 146, 148, 149, 151, 152, 154, 155, 157, 158, 160, 161, 163, 164,
166, 167, 169, 170, 172, 173, 175, 176, 178, 179, 181, 182, 184, 185, 187, 188,
190, 191, 193, 194, 196, 197, 199, 200, 202, 203, 205, 206, 208, 209, 211, 212,
214, 215, 217, 218, 220, 221, 223, 224, 226, 227, 229, 230, 232, 233, 235, 236,
238, 239, 241, 242, 244, 245, 247, 248, 250, 251, 253, 254, 256, 257, 259, 260,
262, 263, 265, 266, 268, 269, 271, 272, 274, 275, 277, 278, 280, 281, 283, 284,
286, 287, 289, 290, 292, 293, 295, 296, 298, 299, 301, 302, 304, 305, 307, 308,
310, 311, 313, 314, 316, 317, 319, 320, 322, 323, 325, 326, 328, 329, 331, 332,
334, 335, 337, 338, 340, 341, 343, 344, 346, 347, 349, 350, 352, 353, 355, 356,
358, 359, 361, 362, 364, 365, 367, 368, 370, 371, 373, 374, 376, 377, 379, 380,
382, 383, 385, 386, 388, 389, 391, 392, 394, 395, 397, 398, 400, 401, 403, 404,
406, 407, 409, 410, 412, 413, 415, 416, 418, 419, 421, 422, 424, 425, 427, 428,
430, 431, 433, 434, 436, 437, 439, 440, 442, 443, 445, 446, 448, 449, 451, 452,
454, 455, 457, 458, 460, 461, 463, 464, 466, 467, 469, 470, 472, 473, 475, 476,

```

478, 479, 481, 482, 484, 485, 487, 488, 490, 491, 493, 494, 496, 497, 499, 500,
 502, 503, 505, 506, 508, 509, 511, 512, 514, 515, 517, 518, 520, 521, 523, 524,
 526, 527, 529, 530, 532, 533, 535, 536, 538, 539, 541, 542, 544, 545, 547, 548,
 550, 551, 553, 554, 556, 557, 559, 560, 562, 563, 565, 566, 568, 569, 571, 572,
 574, 575, 577, 578, 580, 581, 583, 584, 586, 587, 589, 590, 592, 593, 595, 596,
 598, 599, 601, 602, 604, 605, 607, 608, 610, 611, 613, 614, 616, 617, 619, 620,
 622, 623, 625, 626, 628, 629, 631, 632, 634, 635, 637, 638, 640, 641, 643, 644,
 646, 647, 649, 650, 652, 653, 655, 656, 658, 659, 661, 662, 664, 665, 667, 668,
 670, 671, 673, 674, 676, 677, 679, 680, 682, 683, 685, 686, 688, 689]


```
[61]: #---calculate misclassification error for each k---#
MSE=[1-x for x in cv_scores]

#---determining best k (min.MSE)---#
optimal_k=ks[MSE.index(min(MSE))]
print(f"The optimal number of neighbors is {optimal_k}")
print('This is a simulation where I use all features ')

#---plot misclassification error vs k---#
plt.plot(ks,MSE)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()
```

The optimal number of neighbors is 17

This is a simulation where I use all features



2 LAB5 KNN Classification on Diabetes Part4

For the KNN algorithm with the optimal K in part 4, calculate the accuracy, the precision and the recall. Also graph the ROC and estimate AUC. Use part 1 as a guide. Compare the KNN and the LR algorithms. Which one performs better as a classifier for diabetes prediction. Include a table that compares the accuracy metrics of the two algorithms

```
[23]: X=df_array[:, :8]          ##### select All
      y=df.c # k = target (classes)
```

2.0.1 2. Evaluation procedure 2 - Train/test split¶

1. Split the dataset into two pieces: a training set and a testing set.
2. Train the model on the training set.
3. Test the model on the testing set, and evaluate how well we did.

```
[24]: # print the shapes of X and y
      # X is our features matrix with 150 x 4 dimension
      print(X.shape)
      # y is our response vector with 150 x 1 dimension
      print(y.shape)
```

```
(767, 8)
(767,)
```

```
[25]: # STEP 1: split X and y into training and testing sets
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(
          X,
          y,
          test_size=0.4,
          random_state=2)
```

```
X_train = train_set
X_test = test_set
y_train = train_labels
y_test = test_labels
```

- test_size=0.4 40% of observations to test set
60% of observations to training set
- data is randomly assigned unless you use random_state hyperparameter If you use random_state=4
Your data will be split exactly the same way

```
[26]: # print the shapes of the new X objects
print(X_train.shape)
print(X_test.shape)
```

```
(460, 8)
(307, 8)
```

```
[27]: # print the shapes of the new y objects
print(y_train.shape)
print(y_test.shape)
```

```
(460,)
(307,)
```

```
[62]: knn = KNeighborsClassifier(n_neighbors=17)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("--- Accuracy ---- k = 5")
print(metrics.accuracy_score(y_test, y_pred))
```

```
--- Accuracy ---- k = 5
0.7361563517915309
```

```
[63]: knn = KNeighborsClassifier(n_neighbors=30) # choose k
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("--- Accuracy ---- k = 30")

print(metrics.accuracy_score(y_test, y_pred))
```

```
--- Accuracy ---- k = 30
0.7719869706840391
```

```
[ ]:
```

```
[30]: # try K=1 through K=n and record testing accuracy
k_range = range(1, 40)

# We can create Python dictionary using [] or dict()
scores = {}

# We use a loop through the range 1 to something
# We append the scores in the dictionary
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    scores.append(metrics.accuracy_score(y_test, y_pred))
```



```
print(scores)
```


```
[0.6677524429967426, 0.6872964169381107, 0.6644951140065146, 0.7035830618892508,
0.7133550488599348, 0.7035830618892508, 0.7035830618892508, 0.7263843648208469,
0.7328990228013029, 0.7263843648208469, 0.752442996742671, 0.7328990228013029,
0.7296416938110749, 0.7328990228013029, 0.7231270358306189, 0.7361563517915309,
0.7361563517915309, 0.739413680781759, 0.755700325732899, 0.749185667752443,
0.7361563517915309, 0.745928338762215, 0.742671009771987, 0.752442996742671,
0.749185667752443, 0.752442996742671, 0.762214983713355, 0.762214983713355,
0.762214983713355, 0.7719869706840391, 0.749185667752443, 0.749185667752443,
0.745928338762215, 0.745928338762215, 0.745928338762215, 0.758957654723127,
0.758957654723127, 0.755700325732899, 0.745928338762215]
```

```
[31]: # import Matplotlib (scientific plotting library)
import matplotlib.pyplot as plt

# allow plots to appear within the notebook
%matplotlib inline

# plot the relationship between K and testing accuracy
# plt.plot(x_axis, y_axis)
plt.plot(k_range, scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Testing Accuracy')
```

```
[31]: Text(0, 0.5, 'Testing Accuracy')
```



- Training accuracy rises as model complexity increases
- Testing accuracy penalizes models that are too complex or not complex enough
- For KNN models, complexity is determined by the value of K (lower value = more complex)

```
[32]: # get a predicted probabilities and convert into a dataframe
preds_prob = pd.DataFrame(knn.predict_proba(X_test))
```

```
[33]: # assign column names to prediction
preds_prob.columns = ["Malignant", "Benign"]
```

```
[34]: # get the predicted class labels
preds = knn.predict(X_test)
preds_class = pd.DataFrame(preds)
preds_class.columns = ["Prediction"]
```

```
[35]: # actual diagnosis
original_result = pd.DataFrame(y_test)
original_result.columns = ["Original Result"]
```

```
[36]: # Merge the three dataframes into one
result = pd.concat([preds_prob, preds_class, original_result], axis =1)
print(result.head())
```

	Malignant	Benign	Prediction	Original Result
0	0.948718	0.051282	0.0	NaN
1	0.948718	0.051282	0.0	1.0
2	0.820513	0.179487	0.0	NaN
3	0.256410	0.743590	1.0	1.0
4	0.974359	0.025641	0.0	NaN

2.0.2 Getting the Confusion Matrix

```
[38]: # generate table of predictions vs actual
print("--- confusion Matrix ---")
print(pd.crosstab(preds, y_test))
```

```
--- confusion Matrix ---
c      0    1
row_0
0      175  52
1       26  54
```

```
[39]: # --- view the confusion matrix ---
print(metrics.confusion_matrix(y_true = y_test,
                               y_pred = preds))
```

```
[[175  26]
 [ 52  54]]
```

```
[40]: # View summary of common classification metrics
print("---- Metrics----")

print(metrics.classification_report(
    y_true = y_test,
    y_pred = preds))
```

```
---- Metrics----
```

	precision	recall	f1-score	support
0	0.77	0.87	0.82	201
1	0.68	0.51	0.58	106
accuracy			0.75	307
macro avg	0.72	0.69	0.70	307
weighted avg	0.74	0.75	0.74	307

2.0.3 Receiving Operating Characteristic (ROC) Curve

```
[41]: # find the predicted probabilities using the test set

probs = knn.predict_proba(X_test)
preds = probs[:,1]

# find the FPR, TPR, and threshold
fpr, tpr, threshold = roc_curve(y_test, preds)
```

```
[42]: print(fpr)
print(tpr)
print(threshold)
```

```
[0.          0.          0.          0.00995025 0.01492537 0.02487562
 0.02985075 0.03482587 0.07462687 0.09452736 0.12935323 0.15920398
 0.17910448 0.26865672 0.30348259 0.35323383 0.39303483 0.43781095
 0.46766169 0.51741294 0.54726368 0.5721393  0.64676617 0.68159204
 0.83084577 0.89552239 0.94527363 1.          ]
[0.          0.00943396 0.05660377 0.10377358 0.13207547 0.1509434
 0.19811321 0.22641509 0.35849057 0.46226415 0.50943396 0.56603774
 0.59433962 0.66981132 0.71698113 0.75471698 0.81132075 0.8490566
 0.87735849 0.90566038 0.93396226 0.95283019 0.95283019 0.96226415]
```

```

0.98113208 0.99056604 1.          1.          ]
[1.76923077 0.76923077 0.74358974 0.69230769 0.66666667 0.64102564
0.61538462 0.58974359 0.56410256 0.53846154 0.51282051 0.48717949
0.46153846 0.41025641 0.38461538 0.35897436 0.33333333 0.30769231
0.28205128 0.25641026 0.23076923 0.20512821 0.17948718 0.15384615
0.1025641  0.07692308 0.05128205 0.02564103]

```

2.0.4 Plotting the ROC and Finding the Area Under the Curve (AUC)

```


[43]: # find the area under the curve
roc_auc = auc(fpr, tpr)

```

```

[44]: plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.plot([0,1], [0,1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel('True Positive Rate (TPR)')
plt.xlabel('False Positive Rate (FPR)')
plt.title('KNN: Receiver Operating Characteristic (ROC)')
plt.legend(loc = 'lower right')
plt.show()

```



3 Report:

Accuracy LR:0.78, vs KNN(k=30): 0.77 #### **Precision**
LR:0.79(0), 0.77(1) vs KNN: 0.63(0), 0.30(1) #### **Recall** LR:0.91(0), 0.54(1) vs KNN: 0.67(0),
0.26(1) #### **AUC value** LR:0.86, vs KNN: 0.45

[]: