


BME_455_Spring_2021_Lab_10

April 29, 2021

```
[1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: df = pd.read_csv("kmeans.csv")
plt.scatter(df['x'],df['y'], c='r', s=18)
```

```
[2]: <matplotlib.collections.PathCollection at 0x7fe18106adc0>
```



```
[3]: #---let k assume a value---
k = 3
```

```
[4]: X = np.array(list(zip(df['x'],df['y'])))
```


```
[5]: #---generate k random points (centroids)---
Cx = np.random.randint(np.min(X[:,0]), np.max(X[:,0]), size = k)
Cy = np.random.randint(np.min(X[:,1]), np.max(X[:,1]), size = k)
```

```
[6]: #---represent the k centroids as a matrix---
C = np.array(list(zip(Cx, Cy)), dtype=np.float64)
print(C)
```

```
[[1. 3.]
 [6. 1.]
 [7. 6.]]
```

```
[7]: #---plot the original points as well as the k centroids---
plt.scatter(df['x'], df['y'], c='r', s=8)
plt.scatter(Cx, Cy, marker='*', c='g', s=160)
plt.xlabel("x")
plt.ylabel("y")
```

```
[7]: Text(0, 0.5, 'y')
```



```
[8]: from copy import deepcopy
```

```
[9]: #---to calculate the distance between two points---
def euclidean_distance(a, b, ax=1):
    return np.linalg.norm(a - b, axis=ax)
```

```

[10]: #---create a matrix of 0 with same dimension as C (centroids)---
C_prev = np.zeros(C.shape)

[11]: #---to store the cluster each point belongs to---
clusters = np.zeros(len(X))


[12]: #---C is the random centroids and C_prev is all 0s---
#---measure the distance between the centroids and C_prev---
distance_differences = euclidean_distance(C, C_prev)

[13]: #---loop as long as there is still a difference in
# distance between the previous and current centroids---
while distance_differences.any() != 0:
    #---assign each value to its closest cluster---
    for i in range(len(X)):
        distances = euclidean_distance(X[i], C)
        #---returns the indices of the minimum values along an axis---
        cluster = np.argmin(distances)
        clusters[i] = cluster
    #---store the prev centroids---
    C_prev = deepcopy(C)
    #---find the new centroids by taking the average value---
    for i in range(k): #---k is the number of clusters---
        #---take all the points in cluster i---
        points = [X[j] for j in range(len(X)) if clusters[j] == i]
        if len(points) != 0:
            C[i] = np.mean(points, axis=0)
    #---find the distances between the old centroids and the new centroids
    distance_differences = euclidean_distance(C, C_prev)

[14]: #---plot the scatter plot---
colors = ['b', 'r', 'y', 'g', 'c', 'm']
for i in range(k):
    points = np.array([X[j] for j in range(len(X)) if clusters[j] == i])
    if len(points) > 0:
        plt.scatter(points[:, 0], points[:, 1], s=10, c=colors[i])
else:
    print("Please regenerate your centroids again.")
    plt.scatter(points[:, 0], points[:, 1], s=10, c=colors[i])
    plt.scatter(C[:, 0], C[:, 1], marker='*', s=100, c='black')

```

Please regenerate your centroids again.



```
[15]: for i, cluster in enumerate(clusters):
      print("Point " + str(X[i]),
            "Cluster " + str(int(cluster)))
```

```
Point [1 1] Cluster 0
Point [2 2] Cluster 0
Point [2 3] Cluster 0
Point [1 4] Cluster 0
Point [3 3] Cluster 0
Point [6 7] Cluster 2
Point [7 8] Cluster 2
Point [6 8] Cluster 2
Point [7 6] Cluster 1
Point [6 9] Cluster 2
Point [2 5] Cluster 0
Point [7 8] Cluster 2
Point [8 9] Cluster 2
Point [7 8] Cluster 2
Point [8 9] Cluster 2
Point [6 7] Cluster 2
Point [7 8] Cluster 2
Point [3 1] Cluster 0
Point [8 4] Cluster 1
Point [8 6] Cluster 1
Point [8 9] Cluster 2
```

```
[16]: print(C)
```

```
[[2.          2.71428571]
 [7.66666667  5.33333333]
 [6.90909091  8.18181818]]
```

**** Using K-Means in Scikit-learn****

```
[17]: #---using sci-kit-learn---
      from sklearn.cluster import KMeans
      k=3
      kmeans = KMeans(n_clusters=k)
```

```
[18]: kmeans = kmeans.fit(X)
```

```
[19]: labels = kmeans.predict(X)
```

```
[20]: centroids = kmeans.cluster_centers_
```

```
[21]: print(labels)
```

```
[1 1 1 1 1 2 2 2 0 2 1 2 2 2 2 2 1 0 0 2]
```


```
[22]: print(centroids)
```

```
[[7.66666667  5.33333333]
 [2.          2.71428571]
 [6.90909091  8.18181818]]
```

Plot the points and centroids on a scatter plot

```
[23]: #---map the labels to colors---
      c = ['b','r','y','g','c','m']
      colors = [c[i] for i in labels]
      plt.scatter(df['x'],df['y'], c=colors, s=18)
      plt.scatter(centroids[:, 0], centroids[:, 1], marker='*', s=100, c='black')
```

```
[23]: <matplotlib.collections.PathCollection at 0x7fe1824fd340>
```



```
[24]: #---making predictions---
cluster = kmeans.predict([[3,4]])[0]
print(c[cluster]) # r
cluster = kmeans.predict([[7,5]])[0]
print(c[cluster]) # y
```

r
b

Calculating the Silhouette Coefficient

```
[25]: from sklearn import metrics
silhouette_samples = metrics.silhouette_samples(X, kmeans.labels_)
print(silhouette_samples)
print("Average of Silhouette Coefficients for k =", k)
print("=====")
print("Silhouette mean:", silhouette_samples.mean())
```

```
[0.67534567 0.73722797 0.73455072 0.66254937 0.6323039 0.29124047
0.65958544 0.56877212 0.30782271 0.58882096 0.51065918 0.65958544
0.58208055 0.65958544 0.58208055 0.29124047 0.65958544 0.60168807
0.52022287 0.42808391 0.58208055]
Average of Silhouette Coefficients for k = 3
=====
Silhouette mean: 0.5683386570049157
```

```
[26]: print("Silhouette mean:", metrics.silhouette_score(X, kmeans.labels_))
```

Silhouette mean: 0.5683386570049157

Finding the Optimal K

```
[27]: silhouette_avgs = []  
min_k = 2
```

```
[28]: #---try k from 2 to maximum number of labels---  
for k in range(min_k, len(X)):  
    kmean = KMeans(n_clusters=k).fit(X)  
    score = metrics.silhouette_score(X, kmean.labels_)  
    print("Silhouette Coefficients for k =", k, "is", score)  
    silhouette_avgs.append(score)  
f, ax = plt.subplots(figsize=(7, 5))  
ax.plot(range(min_k, len(X)), silhouette_avgs)  
plt.xlabel("Number of clusters")  
plt.ylabel("Silhouette Coefficients")  
#---the optimal k is the one with the highest average silhouette---  
Optimal_K = silhouette_avgs.index(max(silhouette_avgs)) + min_k  
print("Optimal K is ", Optimal_K)
```

```
Silhouette Coefficients for k = 2 is 0.704998665863519  
Silhouette Coefficients for k = 3 is 0.5683386570049157  
Silhouette Coefficients for k = 4 is 0.4643763365584561  
Silhouette Coefficients for k = 5 is 0.48164708359899294  
Silhouette Coefficients for k = 6 is 0.4575232132438315  
Silhouette Coefficients for k = 7 is 0.44821311022689675  
Silhouette Coefficients for k = 8 is 0.5060017736921864  
Silhouette Coefficients for k = 9 is 0.49578816225194444  
Silhouette Coefficients for k = 10 is 0.5588483560001315  
Silhouette Coefficients for k = 11 is 0.52007489882161  
Silhouette Coefficients for k = 12 is 0.4942228407828742  
Silhouette Coefficients for k = 13 is 0.4802755446489002  
Silhouette Coefficients for k = 14 is 0.4425187247054025  
Silhouette Coefficients for k = 15 is 0.42857142857142855  
Silhouette Coefficients for k = 16 is 0.42857142857142855
```

```
<ipython-input-28-2cb04cdee8d7>:3: ConvergenceWarning: Number of distinct  
clusters (15) found smaller than n_clusters (16). Possibly due to duplicate  
points in X.
```

```
kmean = KMeans(n_clusters=k).fit(X)
```

```
<ipython-input-28-2cb04cdee8d7>:3: ConvergenceWarning: Number of distinct  
clusters (15) found smaller than n_clusters (17). Possibly due to duplicate  
points in X.
```

```
kmean = KMeans(n_clusters=k).fit(X)
```

```
Silhouette Coefficients for k = 17 is 0.42857142857142855
```

```
<ipython-input-28-2cb04cdee8d7>:3: ConvergenceWarning: Number of distinct  
clusters (15) found smaller than n_clusters (18). Possibly due to duplicate
```

points in X.

```
kmean = KMeans(n_clusters=k).fit(X)
```

Silhouette Coefficients for k = 18 is 0.42857142857142855

<ipython-input-28-2cb04cdee8d7>:3: ConvergenceWarning: Number of distinct clusters (15) found smaller than n_clusters (19). Possibly due to duplicate points in X.

```
kmean = KMeans(n_clusters=k).fit(X)
```


Silhouette Coefficients for k = 19 is 0.42857142857142855

Silhouette Coefficients for k = 20 is 0.42857142857142855

Optimal K is 2

<ipython-input-28-2cb04cdee8d7>:3: ConvergenceWarning: Number of distinct clusters (15) found smaller than n_clusters (20). Possibly due to duplicate points in X.

```
kmean = KMeans(n_clusters=k).fit(X)
```



```
[29]: %matplotlib inline
import numpy as np
import pandas as pd
df = pd.read_csv("BMX_G.csv")
```



```
[30]: print(df.shape)
```

```
(9338, 27)
```

```
[31]: df.isnull().sum()
```


```
[31]: Unnamed: 0      0
      seqn          0
      bmdstats      0
      bmxwt          95
      bmiwt        8959
      bmxrecum      8259
      bmirecum      9307
      bmxhead       9102
      bmihead       9338
      bmxht         723
      bmiht         9070
      bmxbmi         736
      bmdbmic        5983
      bmxleg        2383
      bmileg        8984
      bmxarml        512
      bmiarml       8969
      bmxarmc        512
      bmiarmc       8965
      bmxwaist       1134
      bmiwaist       8882
      bmxsad1        2543
      bmxsad2        2543
      bmxsad3        8940
      bmxsad4        8940
      bmdavsad       2543
      bmdsadcml      8853
      dtype: int64
```

```
[32]: df = df.dropna(subset=['bmxleg','bmxwaist']) # remove rows with NaNs
      print(df.shape)
```

```
(6899, 27)
```

```
[33]: import matplotlib.pyplot as plt
      plt.scatter(df['bmxleg'],df['bmxwaist'], c='r', s=2)
      plt.xlabel("Upper leg Length (cm)")
      plt.ylabel("Waist Circumference (cm)")
```

```
[33]: Text(0, 0.5, 'Waist Circumference (cm)')
```




Clustering Using K-Means

```
[34]: #---using sci-kit-learn---
      from sklearn.cluster import KMeans
```

```
[35]: #---using sci-kit-learn---
      from sklearn.cluster import KMeans
      k = 2
      X = np.array(list(zip(df['bmxleg'], df['bmxwaist'])))
      kmeans = KMeans(n_clusters=k)
      kmeans = kmeans.fit(X)
      labels = kmeans.predict(X)
      centroids = kmeans.cluster_centers_
      #---map the labels to colors---
      c = ['b', 'r', 'y', 'g', 'c', 'm']
      colors = [c[i] for i in labels]
      plt.scatter(df['bmxleg'], df['bmxwaist'], c=colors, s=2)
      plt.scatter(centroids[:, 0], centroids[:, 1], marker='*', s=100, c='black')
```

```
[35]: <matplotlib.collections.PathCollection at 0x7fe18318e4c0>
```



```
[36]: print(centroids)
```

```
[[ 38.82253829 108.01287902]
 [ 37.65959459  77.92516216]]
```

Finding the Optimal Size Classes

```
[37]: from sklearn import metrics
silhouette_avgs = []
min_k = 2
#---try k from 2 to maximum number of labels---
for k in range(min_k, 10):
    kmean = KMeans(n_clusters=k).fit(X)
    score = metrics.silhouette_score(X, kmean.labels_)
    print("Silhouette Coefficients for k =", k, "is", score)
    silhouette_avgs.append(score)
#---the optimal k is the one with the highest average silhouette---
Optimal_K = silhouette_avgs.index(max(silhouette_avgs)) + min_k
print("Optimal K is", Optimal_K)
```

```
Silhouette Coefficients for k = 2 is 0.5165601620046447
Silhouette Coefficients for k = 3 is 0.47123155436589514
Silhouette Coefficients for k = 4 is 0.4359623202434665
Silhouette Coefficients for k = 5 is 0.41917232165890106
Silhouette Coefficients for k = 6 is 0.3937070096308068
Silhouette Coefficients for k = 7 is 0.37701689327826704
Silhouette Coefficients for k = 8 is 0.3574365786281303
```

Silhouette Coefficients for $k = 9$ is 0.3410434517384024
Optimal K is 2

```
[38]: #trying with k=4
k = 4
X = np.array(list(zip(df['bmxleg'],df['bmxwaist'])))
kmeans = KMeans(n_clusters=k)
kmeans = kmeans.fit(X)
labels = kmeans.predict(X)
centroids = kmeans.cluster_centers_
#---map the labels to colors---
c = ['b','r','y','g','c','m']
colors = [c[i] for i in labels]
plt.scatter(df['bmxleg'],df['bmxwaist'], c=colors, s=2)
plt.scatter(centroids[:, 0], centroids[:, 1], marker='*', s=100, c='black')
```

[38]: <matplotlib.collections.PathCollection at 0x7fe1831cfb20>

